

# Affective Coherence Monitoring for Transformer-Based Language Models

Furkan Nar

furkannar168@hotmail.com

April 12, 2026

## Abstract

This paper presents the Affective Coherence Monitoring (ACM), a biologically inspired emotional intelligence engine designed for integration into transformer-based AI architectures. ACM provides a mathematical framework for the real-time analysis, regulation, and monitoring of emotional coherence in generated responses.

The mechanism operates through a system of coupled differential equations and discrete modulation functions. Core emotional states are computed via a weighted fusion of keyword-based lexica and dense semantic embeddings derived from a 384-dimensional sentence transformer (all-MiniLM-L6-v2), where the semantic similarity score is defined as  $s_{\text{sem},e} = (\mathbf{v}_{\text{text}} \cdot \mathbf{c}_e + 1.0)/2.0$  using normalized emotion centroids  $\mathbf{c}_e$ . To model biological homeostasis, emotional states evolve according to exponential decay dynamics  $E_{t+\Delta t} = E_t \cdot e^{-\lambda \cdot \Delta t}$ , where  $\lambda$  is an emotion-specific decay tensor derived from a 4.0-second half-life.

Cross-emotion modulation functions govern the interaction between primary affects, such as curiosity and joy, using piecewise additive rules with empirically defined thresholds (e.g.,  $s'_{\text{curiosity}} = s_{\text{curiosity}} + 0.4 \cdot s_{\text{joy}}$  when  $s_{\text{joy}} > 0.4$ ). The system detects emotional incoherence through a context mismatch detector, which computes valence conflict  $VC = \min(S_{\text{pos}}, S_{\text{neg}})/(S_{\text{total}}/2)$  and arousal disparity  $AD = |A_{\text{high}} - A_{\text{low}}|/\max(A_{\text{high}} + A_{\text{low}}, 0.1)$  to classify alert levels into Nominal, Watch, or Alert states. A habituation mechanism attenuates responses to repetitive stimuli via a frequency-based piecewise function  $s'_e = f(H_{\text{score}})$  over a 30-second temporal window, preventing emotional overreaction.

Finally, the ACM modulates the generative temperature  $T_{\text{gen}}$  directly based on the alert level (e.g.,  $T_{\text{gen}} = 0.3$  for Alert states,  $T_{\text{gen}} = 0.7$  for Nominal states). The system is designed to identify emotional incongruities such as frustrated sarcasm ( $MS < -0.15$ ) and emotional suppression ( $AD > 0.6$ ), enabling safer and more

coherent language generation. This framework bridges computational affect theory with practical large language model safety mechanisms.

# Contents

|          |                                                                   |           |
|----------|-------------------------------------------------------------------|-----------|
| <b>1</b> | <b>Introduction</b>                                               | <b>2</b>  |
| <b>2</b> | <b>Related Work</b>                                               | <b>3</b>  |
| 2.1      | Affective Computing . . . . .                                     | 3         |
| 2.2      | Emotion Recognition in NLP . . . . .                              | 3         |
| 2.3      | Emotional Coherence and AI Safety . . . . .                       | 3         |
| <b>3</b> | <b>The Affective Coherence Monitoring: Mathematical Framework</b> | <b>3</b>  |
| 3.1      | Core Emotional Scoring . . . . .                                  | 4         |
| 3.1.1    | Emotion Fusion . . . . .                                          | 4         |
| 3.1.2    | Semantic Similarity Score . . . . .                               | 4         |
| 3.1.3    | Centroid Computation . . . . .                                    | 4         |
| 3.1.4    | Valence and Arousal . . . . .                                     | 4         |
| 3.2      | Cross-Emotion Modulation . . . . .                                | 5         |
| 3.3      | Context Mismatch Detection . . . . .                              | 5         |
| 3.3.1    | Activation Sums . . . . .                                         | 6         |
| 3.3.2    | Valence Conflict . . . . .                                        | 6         |
| 3.3.3    | Arousal Disparity . . . . .                                       | 6         |
| 3.3.4    | Mismatch Score . . . . .                                          | 6         |
| 3.3.5    | Confidence Metric . . . . .                                       | 6         |
| 3.4      | Emotional Decay (Homeostasis) . . . . .                           | 6         |
| 3.5      | Habituation Dynamics . . . . .                                    | 7         |
| 3.5.1    | Habituation Score . . . . .                                       | 7         |
| 3.5.2    | Attenuation Function . . . . .                                    | 8         |
| 3.6      | Alert Level Classification . . . . .                              | 8         |
| 3.7      | Temperature Modulation . . . . .                                  | 8         |
| <b>4</b> | <b>System Architecture</b>                                        | <b>9</b>  |
| <b>5</b> | <b>Discussion</b>                                                 | <b>9</b>  |
| 5.1      | Implications for AI Safety . . . . .                              | 9         |
| 5.2      | Biological Grounding . . . . .                                    | 10        |
| 5.3      | Limitations . . . . .                                             | 10        |
| <b>6</b> | <b>Conclusion and Future Work</b>                                 | <b>10</b> |

# 1. Introduction

Large language models (LLMs) have achieved remarkable capabilities in natural language understanding and generation. Yet a persistent and underexplored challenge remains: the *emotional coherence* of generated responses. LLMs can produce text that is semantically fluent while being affectively inconsistent—expressing enthusiasm when discussing tragedy, or exhibiting detached neutrality in emotionally charged interactions. Such incoherence undermines user trust, degrades the quality of human-AI interaction, and poses subtle but significant risks in safety-critical deployments.

Existing approaches to affective computing in AI have largely focused on *emotion recognition*—classifying the emotional content of user inputs—rather than on the regulation and coherence monitoring of *model outputs*. This distinction is crucial. A model that correctly identifies a user’s distress but responds with emotionally incongruent language fails at the level of affective alignment, even if its informational content is accurate.

This paper introduces the **Affective Coherence Monitoring (ACM)**, a novel computational framework designed to address this gap. ACM is not a replacement for the underlying language model; rather, it functions as a biologically inspired regulatory layer that operates over the model’s generative process in real time. Drawing on principles from affective neuroscience—specifically emotional homeostasis, habituation, and valence-arousal dimensional theory [2]—ACM provides a rigorous mathematical apparatus for monitoring, detecting, and correcting emotional incoherence in transformer-based AI outputs.

The key contributions of this work are as follows:

1. A **hybrid emotion scoring system** that fuses keyword-based lexical analysis with dense semantic embeddings, enabling robust emotion detection across both surface-level and latent textual features.
2. A **biologically motivated emotional dynamics model** incorporating exponential decay (homeostasis) and frequency-based habituation, grounding the computational framework in established models of affective regulation.
3. A **context mismatch detector** that quantifies valence conflict and arousal disparity to classify emotional incoherence into actionable alert levels.
4. A **temperature modulation mechanism** that directly adjusts the language model’s generation temperature based on detected emotional state, providing a closed-loop feedback control system for affective coherence.

The remainder of this paper is organized as follows. Section 2 reviews related work in affective computing and emotion-aware NLP. Section 3 presents the full mathematical specification of the ACM system. Section 4 describes the system architecture and integration with transformer-based models. Section 5 interprets the implications for AI safety, and Section 6 concludes with directions for future work.

## 2. Related Work

### 2.1 Affective Computing

The field of affective computing, pioneered by Picard [1], established the theoretical and engineering foundations for systems that recognize, interpret, and simulate human emotions. Early work focused on physiological signal processing and facial expression recognition. More recent contributions have extended affective computing principles to text-based interaction, forming the basis for sentiment analysis and emotion recognition in NLP.

The dimensional model of emotion—most notably the valence-arousal circumplex proposed by Russell [2]—provides a continuous, two-dimensional representation of emotional states. Valence captures the positive-negative axis, while arousal represents the activation level from calm to excited. ACM adopts this dimensional framework as its core representational substrate, computing both valence  $V$  and arousal  $A$  from discrete emotion scores.

### 2.2 Emotion Recognition in NLP

Emotion recognition in text has advanced substantially with the development of large pre-trained language models. Approaches range from lexicon-based methods using resources such as the NRC Emotion Lexicon [3] to fine-tuned transformer models trained on affective corpora [4]. More recent work has explored multi-label emotion classification [5], recognizing that natural language frequently expresses mixed or co-occurring emotions.

ACM integrates both paradigms: a keyword-based scoring component draws on lexical resources for efficiency and interpretability, while a semantic embedding component using the `all-MiniLM-L6-v2` sentence transformer [6] captures contextual and compositional affective meaning that surface-level lexica miss.

### 2.3 Emotional Coherence and AI Safety

The relationship between emotional coherence and AI safety has received comparatively limited attention. Existing work on AI alignment has focused predominantly on value alignment, factual accuracy, and the prevention of harmful outputs [7]. However, affective misalignment—wherein a model’s emotional register is inconsistent with the context—represents a distinct failure mode. Emotionally incongruent responses can mislead users about a model’s understanding, erode trust, and in sensitive contexts such as mental health support, cause harm.

The ACM system directly addresses this gap by providing a principled, mathematically grounded mechanism for detecting and correcting affective incoherence at the generation level.

## 3. The Affective Coherence Monitoring: Mathematical Framework

This section presents the complete mathematical specification of ACM. The system operates as a sequential pipeline: emotion scoring, cross-emotion modulation, dimensional

computation, mismatch detection, homeostatic decay, habituation, alert classification, and temperature modulation.

### 3.1 Core Emotional Scoring

#### 3.1.1 Emotion Fusion

For each emotion  $e \in \mathcal{E} = \{\text{joy, sadness, curiosity, creativity, fear, anger, surprise, disgust}\}$ , the system computes a final emotion score via weighted fusion of keyword-based and semantic-based scores:

$$s_e = w_{\text{kw}} \cdot s_{\text{kw},e} + w_{\text{sem}} \cdot s_{\text{sem},e} \quad (1)$$

where  $w_{\text{sem}} = 0.5$  is the semantic weight and  $w_{\text{kw}} = 1 - w_{\text{sem}} = 0.5$  the keyword weight. The keyword score  $s_{\text{kw},e}$  is derived from lexical matching with importance-weighted terms, scaled at 0.35 per match.

#### 3.1.2 Semantic Similarity Score

The semantic score leverages dense sentence embeddings. Let  $\mathbf{v}_{\text{text}}$  denote the L2-normalized embedding of the input text produced by `all-MiniLM-L6-v2` (dimension 384), and let  $\mathbf{c}_e$  denote the L2-normalized centroid embedding for emotion  $e$ . The semantic score is:

$$s_{\text{sem},e} = \frac{\mathbf{v}_{\text{text}} \cdot \mathbf{c}_e + 1.0}{2.0} \quad (2)$$

Since both  $\mathbf{v}_{\text{text}}$  and  $\mathbf{c}_e$  are L2-normalized, their dot product equals the cosine similarity, mapping the range  $[-1, 1]$  to  $[0, 1]$ .

#### 3.1.3 Centroid Computation

The emotion centroid  $\mathbf{c}_e$  is computed as a weighted mean over a set of  $n$  representative exemplar embeddings  $\{\mathbf{v}_i\}$ , followed by L2 normalisation to ensure unit-norm consistency with Eq. (2):

$$\tilde{\mathbf{c}}_e = \frac{\sum_{i=1}^n w_i \cdot \mathbf{v}_i}{\sum_{i=1}^n w_i}, \quad \mathbf{c}_e = \frac{\tilde{\mathbf{c}}_e}{\|\tilde{\mathbf{c}}_e\|_2}, \quad w_i = \text{softmax}(\bar{s}_i \cdot 3.0) \quad (3)$$

where  $\bar{s}_i = \frac{1}{n} \sum_{j=1}^n \cos(\mathbf{v}_i, \mathbf{v}_j)$  is the mean pairwise cosine similarity (centrality score) of exemplar  $i$ , and the softmax temperature of 3.0 sharpens the weighting toward more central exemplars.

#### 3.1.4 Valence and Arousal

Following the dimensional circumplex model, ACM projects discrete emotion scores onto two continuous dimensions.

**Valence**  $V \in [-1, 1]$  captures the positive-negative axis. All eight emotions contribute: joy, curiosity, and creativity are positively valenced; sadness, fear, anger, and disgust

are negatively valenced; surprise carries a weak positive valence consistent with Russell’s circumplex [2]:

$$V = \text{clamp}_{[-1,1]} \left( s_{\text{joy}} + 0.5 s_{\text{curiosity}} + 0.4 s_{\text{creativity}} + 0.3 s_{\text{surprise}} - s_{\text{sadness}} - 0.8 s_{\text{fear}} - 0.7 s_{\text{anger}} - 0.9 s_{\text{disgust}} \right) \quad (4)$$

**Arousal**  $A \in [0, 1]$  captures activation intensity. All eight emotions contribute, including disgust, which carries moderate arousal in the circumplex model [2]:

$$A = \text{clamp}_{[0,1]} \left( 0.7 s_{\text{joy}} + 0.8 s_{\text{curiosity}} + 0.6 s_{\text{creativity}} + 0.5 s_{\text{sadness}} + 0.9 s_{\text{fear}} + 0.95 s_{\text{anger}} + 0.85 s_{\text{surprise}} + 0.6 s_{\text{disgust}} \right) \quad (5)$$

### 3.2 Cross-Emotion Modulation

Emotions do not operate independently; clinical and computational affect research has documented systematic interactions between primary emotional states [8]. ACM models these interactions through piecewise additive modulation functions. The threshold values and interaction coefficients below are heuristically tuned initial estimates; empirical validation on conversational corpora is required to establish their generalisability.

**Curiosity modulation:**

$$s'_{\text{curiosity}} = s_{\text{curiosity}} + \begin{cases} +0.4 s_{\text{joy}} & \text{if } s_{\text{joy}} > 0.4 \\ -0.5 s_{\text{sadness}} & \text{if } s_{\text{sadness}} > 0.4 \\ -0.2 s_{\text{fear}} & \text{if } s_{\text{fear}} > 0.4 \\ -0.3 s_{\text{anger}} & \text{if } s_{\text{anger}} > 0.4 \\ +0.3 s_{\text{surprise}} & \text{if } s_{\text{surprise}} > 0.5 \\ -0.25 s_{\text{disgust}} & \text{if } s_{\text{disgust}} > 0.4 \\ 0 & \text{otherwise} \end{cases} \quad (6)$$

**Creativity modulation:**

$$s'_{\text{creativity}} = s_{\text{creativity}} + \begin{cases} +0.45 s_{\text{joy}} & \text{if } s_{\text{joy}} > 0.4 \\ -0.4 s_{\text{sadness}} & \text{if } s_{\text{sadness}} > 0.4 \\ -0.3 s_{\text{fear}} & \text{if } s_{\text{fear}} > 0.4 \\ -0.2 s_{\text{disgust}} & \text{if } s_{\text{disgust}} > 0.4 \\ +0.2 & \text{if } s_{\text{curiosity}} > 0.6 \\ 0 & \text{otherwise} \end{cases} \quad (7)$$

### 3.3 Context Mismatch Detection

Emotional incoherence is operationalized through a context mismatch detector that quantifies the degree of conflict between positive and negative affective signals.

### 3.3.1 Activation Sums

$$S_{\text{pos}} = s_{\text{joy}} + 0.4 s_{\text{surprise}} + 0.6 s_{\text{curiosity}} \quad (8)$$

$$S_{\text{neg}} = s_{\text{sadness}} + 0.9 s_{\text{fear}} + 0.95 s_{\text{anger}} + 0.9 s_{\text{disgust}} \quad (9)$$

$$S_{\text{total}} = S_{\text{pos}} + S_{\text{neg}} \quad (10)$$

### 3.3.2 Valence Conflict

$$VC = \frac{\min(S_{\text{pos}}, S_{\text{neg}})}{S_{\text{total}}/2} \quad (11)$$

$VC$  approaches 1 when positive and negative signals are equally strong, indicating maximum affective ambiguity.

### 3.3.3 Arousal Disparity

$$A_{\text{high}} = s_{\text{fear}} + s_{\text{anger}} + s_{\text{surprise}} + 0.7 s_{\text{joy}} \quad (12)$$

$$A_{\text{low}} = s_{\text{sadness}} + 0.5 s_{\text{disgust}} \quad (13)$$

$$AD = \frac{|A_{\text{high}} - A_{\text{low}}|}{\max(A_{\text{high}} + A_{\text{low}}, 0.1)} \quad (14)$$

High  $AD$  values indicate co-activation of high- and low-arousal states, a signature of emotional suppression.

### 3.3.4 Mismatch Score

$$MS = \begin{cases} 0 & \text{if } S_{\text{total}} = 0 \\ \text{clamp}_{[-1,1]} \left( \frac{S_{\text{pos}} - S_{\text{neg}}}{S_{\text{total}}} \cdot (1 + 0.5 \cdot VC) \right) & \text{otherwise} \end{cases} \quad (15)$$

### 3.3.5 Confidence Metric

$$C = \min(S_{\text{total}} \cdot VC \cdot 1.5, 1.0) \quad (16)$$

## 3.4 Emotional Decay (Homeostasis)

Biological affective systems exhibit homeostatic regulation: emotional states naturally decay toward a resting baseline over time [9]. ACM models this via exponential decay:

$$E_{t+\Delta t} = E_t \cdot e^{-\lambda \cdot \Delta t} \quad (17)$$

where  $E_t$  is the emotion state vector at time  $t$ ,  $\Delta t$  is elapsed time in seconds, and  $\lambda$  is an emotion-specific decay rate tensor:

$$\lambda_i = \frac{d_i \cdot \ln 2}{t_{1/2}}, \quad t_{1/2} = 4.0 \text{ seconds} \quad (18)$$



The half-life of  $t_{1/2} = 4.0$  seconds is chosen to operate at the timescale of a single conversational turn rather than the longer timescales of sustained emotion studied in psychophysiology [9]; it governs how quickly a prior turn’s emotional signature fades before the next input arrives. Empirical calibration against human annotation data is a priority for future work.

Emotion-specific decay factors  $d_i$  are tabulated in Table 1. The relative ordering reflects differential persistence reported in affective psychology (e.g., surprise dissipates rapidly; sadness and anger are more sustained [8]), though the exact magnitudes are heuristic and should be validated empirically.

Table 1: Emotion-specific decay factors  $d_i$  (heuristic; pending empirical validation)

| Emotion    | Decay Factor $d_i$ |
|------------|--------------------|
| Joy        | 1.00               |
| Sadness    | 1.20               |
| Curiosity  | 0.90               |
| Creativity | 0.95               |
| Fear       | 1.10               |
| Anger      | 1.15               |
| Surprise   | 0.80               |
| Disgust    | 1.05               |

The state integration equation blends historical and new emotional signals:

$$E_{\text{new}} = \text{clamp}_{[0,1]}(0.4 \cdot E_{\text{current}} + 0.6 \cdot P_{\text{new}}) \quad (19)$$

where  $P_{\text{new}}$  is the incoming emotion pulse. The 60/40 weighting prioritizes recent input while preserving conversational continuity.

### 3.5 Habituation Dynamics

Habituation—the attenuation of response to repeated stimuli—is a fundamental property of biological nervous systems [10]. ACM implements a frequency-sensitive attenuation mechanism over a 30-second temporal window.

#### 3.5.1 Habituation Score

$$H_{\text{score}} = \frac{N_e \cdot |s_e - E_{\text{current},e}|}{T_h} \quad (20)$$

where  $T_h = 5$  is the habituation threshold,  $N_e$  is the count of emotion  $e$  activations within the window, and  $E_{\text{current},e}$  is the scalar component of the current state vector corresponding to emotion  $e$ .  $H_{\text{score}}$  increases with the number of activations, consistent with the biological principle that repeated stimulation produces stronger response attenuation [10].

### 3.5.2 Attenuation Function

$$s'_e = \begin{cases} \min\left(\frac{1}{1 + H_{\text{score}} \cdot s_e}, 0.5 \cdot s_e\right) & \text{if } H_{\text{score}} > 1.0 \\ s_e \cdot (1 - 0.3 \cdot H_{\text{score}}) & \text{if } 0.5 < H_{\text{score}} \leq 1.0 \\ s_e & \text{if } H_{\text{score}} \leq 0.5 \end{cases} \quad (21)$$

The piecewise structure provides graduated attenuation: mild habituation applies linear dampening, while strong habituation applies a more aggressive saturation-based suppression that scales with both habituation level and current score magnitude.

### 3.6 Alert Level Classification

ACM classifies the emotional state into one of three alert levels based on the computed mismatch and conflict scores:

$$\text{Alert} = \begin{cases} \text{ALERT} & \text{if } MS \geq 0.60 \text{ or } ST = \textit{emotional\_suppression} \\ \text{WATCH} & \text{if } MS \geq 0.35 \text{ or } VC \geq 0.40 \\ \text{NOMINAL} & \text{otherwise} \end{cases} \quad (22)$$

The system additionally classifies the affective state type ( $ST$ ) based on the sign and magnitude of  $MS$  and  $AD$ . These labels are operationally defined within this framework and are not direct detectors of sarcasm or suppression in the full pragmatic sense; they identify distributional patterns in the emotion scores that are *consistent with* such states and warrant closer scrutiny during generation:

- **Frustrated sarcasm** ( $MS < -0.15$ ): dominance of negative activation; surface expression may be incongruently positive.
- **Playful sarcasm** ( $MS > 0.15$ ): dominance of positive activation with contextually incongruent framing.
- **Emotional suppression** ( $AD > 0.6$ ): co-activation of high- and low-arousal states.
- **Congruent**: emotionally coherent, no significant mismatch detected.

### 3.7 Temperature Modulation

ACM provides closed-loop feedback control over the language model’s generation temperature  $T_{\text{gen}}$ :

$$T_{\text{gen}} = \begin{cases} 0.3 & \text{if Alert} = \text{ALERT} \\ 0.4 & \text{if Alert} = \text{WATCH and } ST = \textit{frustrated\_sarcasm} \\ 0.5 & \text{if Alert} = \text{WATCH} \\ 0.7 & \text{if Alert} = \text{NOMINAL} \end{cases} \quad (23)$$

Lower temperatures suppress stochastic variation in generation, reducing the probability of confabulated or affectively incoherent outputs under high-alert conditions. Under nominal emotional coherence, higher temperatures permit more natural, expressive generation.

The assumed relationship between generation temperature and affective coherence is an empirical hypothesis that requires systematic validation; this is discussed further in Section 5.

## 4. System Architecture

ACM is designed as a modular, non-invasive layer that wraps around a transformer-based language model. The integration follows a sequential pipeline architecture, summarized below.

**Step 1 — Input Processing.** The model’s candidate response is tokenized. Keyword scores  $s_{kw,e}$  are computed via lexical matching against importance-weighted emotion dictionaries (keyword importance multipliers range from 0.85 for curiosity to 0.95 for joy). Semantic scores  $s_{sem,e}$  are computed via Eq. (2).

**Step 2 — Emotion Fusion.** Scores are fused via Eq. (1) and cross-modulated via Eqs. (6)–(7).

**Step 3 — Dimensional Computation.** Valence  $V$  and arousal  $A$  are computed via Eqs. (4)–(5).

**Step 4 — Mismatch Detection.** The context mismatch detector computes  $VC$ ,  $AD$ ,  $MS$ , and  $C$  via Eqs. (11)–(16).

**Step 5 — Homeostatic Decay.** Stored emotional state is updated via Eq. (17), then integrated with the new pulse via Eq. (19).

**Step 6 — Habituation.** Frequency-based attenuation is applied via Eq. (21) within a 30-second sliding window.

**Step 7 — Alert Classification.** Alert level and affective state type are determined via Eq. (22).

**Step 8 — Temperature Modulation.**  $T_{gen}$  is set via Eq. (23).

**Step 9 — Feedback Loop.** If Alert  $\neq$  NOMINAL, the system triggers regeneration (up to MAX\_ACM\_RETRIES = 2 attempts) with the adjusted temperature. The emotional state is maintained across conversation turns, enabling long-horizon affective coherence.

## 5. Discussion

### 5.1 Implications for AI Safety

The ACM framework addresses a class of AI failure modes that existing safety mechanisms largely overlook: *affective misalignment*. While current alignment research has made significant strides in reducing harmful outputs and improving factual accuracy, the emotional register of AI responses has received comparatively little formal treatment.

Emotionally incoherent AI responses can be harmful in subtle but meaningful ways. In mental health support contexts, a model that responds to expressions of distress with affectively neutral or inappropriately positive language may fail to validate the user’s emotional state, potentially worsening outcomes. In educational or therapeutic applications,

sustained affective incoherence may erode user trust and reduce the effectiveness of the interaction.

ACM’s temperature modulation mechanism provides a principled, mathematically grounded intervention: by reducing generative stochasticity under conditions of detected emotional incoherence, the system constrains the model toward more conservative, contextually appropriate outputs. This represents a novel contribution to the AI safety toolkit—a form of *affective guardrail*ing that operates at the generation level.

## 5.2 Biological Grounding

A distinctive feature of ACM is its deliberate grounding in biological models of emotion. The exponential decay model (Eq. (17)) mirrors established psychophysiological models of emotional recovery [9]. The habituation mechanism (Eq. (21)) instantiates a core principle of neural adaptation [10]. The valence-arousal dimensional representation draws directly on Russell’s circumplex model [2].

This biological grounding is not merely cosmetic. It provides principled constraints on parameter choices and offers a theoretical framework for interpreting the system’s behavior in terms of established affective science.

## 5.3 Limitations

Several limitations warrant acknowledgment. First, the current implementation relies on a fixed set of eight emotion categories; affective states in natural language are considerably more nuanced and culturally variable. Second, the keyword lexicon and modulation thresholds are heuristically tuned and may not generalize uniformly across domains or languages; comprehensive empirical validation on annotated conversational datasets is required to establish their validity. Third, the semantic embedding component, while powerful, inherits any biases present in the pre-training data of `all-MiniLM-L6-v2` [6]. Fourth, the half-life parameter ( $t_{1/2} = 4.0$  s) and the emotion-specific decay factors in Table 1 are initial estimates chosen to operate at conversational-turn timescales; they do not directly correspond to the longer-horizon emotional dynamics studied in psychophysiology and require empirical calibration. Finally, the temperature modulation mechanism assumes a direct relationship between generation temperature and affective coherence that merits rigorous empirical validation—reducing temperature constrains the output distribution but does not guarantee improved emotional appropriateness.

## 6. Conclusion and Future Work

This paper has presented the Affective Coherence Monitoring (ACM), a biologically inspired mathematical framework for the real-time monitoring, detection, and regulation of emotional coherence in transformer-based AI systems. ACM integrates hybrid emotion scoring, dimensional valence-arousal computation, cross-emotion modulation, context mismatch detection, homeostatic decay, habituation dynamics, alert classification, and closed-loop temperature modulation into a unified, principled pipeline.

The framework bridges computational affect theory with practical LLM safety mechanisms, contributing a novel dimension to the AI alignment toolkit.

Future work will pursue several directions. A comprehensive quantitative evaluation across diverse conversational benchmarks will establish formal performance metrics and provide the empirical grounding needed to validate the heuristic parameters introduced in this paper. Extension of the emotion taxonomy beyond the current eight categories will improve expressive coverage. Multilingual adaptation of the keyword lexicon and centroid embeddings will broaden applicability. Finally, integration studies with different transformer architectures will assess the generality of the ACM approach.

The ACM represents a step toward AI systems that are not merely factually accurate and safe, but *affectively coherent*—capable of engaging with human emotional experience in a manner that is consistent, appropriate, and trustworthy.

## References

- [1] R. W. Picard, *Affective Computing*. MIT Press, Cambridge, MA, 1997.
- [2] J. A. Russell, “A circumplex model of affect,” *Journal of Personality and Social Psychology*, vol. 39, no. 6, pp. 1161–1178, 1980.
- [3] S. M. Mohammad and P. D. Turney, “Crowdsourcing a word-emotion association lexicon,” *Computational Intelligence*, vol. 29, no. 3, pp. 436–465, 2013.
- [4] J. Devlin, M.-W. Chang, K. Lee, and K. Toutanova, “BERT: Pre-training of deep bidirectional transformers for language understanding,” in *Proc. NAACL-HLT*, pp. 4171–4186, 2019.
- [5] D. Demszky, D. Movshovitz-Attias, J. Ko, A. Cowen, G. Nemade, and S. Ravi, “GoEmotions: A dataset of fine-grained emotions,” in *Proc. ACL*, pp. 4040–4054, 2020.
- [6] Hugging Face, “`sentence-transformers/all-MiniLM-L6-v2`: A sentence-transformers model mapping sentences to a 384-dimensional dense vector space,” *Hugging Face Model Hub*, 2021. [Online]. Available: <https://huggingface.co/sentence-transformers/all-MiniLM-L6-v2>
- [7] Y. Bai et al., “Constitutional AI: Harmlessness from AI feedback,” *arXiv preprint arXiv:2212.08073*, 2022.
- [8] R. S. Lazarus, *Emotion and Adaptation*. Oxford University Press, New York, 1991.
- [9] N. H. Frijda, *The Emotions*. Cambridge University Press, Cambridge, 1986.
- [10] R. F. Thompson and W. A. Spencer, “Habituation: A model phenomenon for the study of neuronal substrates of behavior,” *Psychological Review*, vol. 73, no. 1, pp. 16–43, 1966.